

# Emotional intelligence of Large Language Models

Journal of Pacific Rim Psychology

Volume 17: 1–12

© The Author(s) 2023


Article reuse guidelines:

sagepub.com/journals-permissions

DOI: 10.1177/18344909231213958

journals.sagepub.com/home/pac



Xuena Wang<sup>1</sup> , Xueting Li<sup>2</sup>, Zi Yin<sup>1</sup>, Yue Wu<sup>1</sup> and Jia Liu<sup>1</sup>

## Abstract

Large Language Models (LLMs) have demonstrated remarkable abilities across numerous disciplines, primarily assessed through tasks in language generation, knowledge utilization, and complex reasoning. However, their alignment with human emotions and values, which is critical for real-world applications, has not been systematically evaluated. Here, we assessed LLMs' Emotional Intelligence (EI), encompassing emotion recognition, interpretation, and understanding, which is necessary for effective communication and social interactions. Specifically, we first developed a novel psychometric assessment focusing on Emotion Understanding (EU), a core component of EI. This test is an objective, performance-driven, and text-based evaluation, which requires evaluating complex emotions in realistic scenarios, providing a consistent assessment for both human and LLM capabilities. With a reference frame constructed from over 500 adults, we tested a variety of mainstream LLMs. Most achieved above-average Emotional Quotient (EQ) scores, with GPT-4 exceeding 89% of human participants with an EQ of 117. Interestingly, a multivariate pattern analysis revealed that some LLMs apparently did not rely on the human-like mechanism to achieve human-level performance, as their representational patterns were qualitatively distinct from humans. In addition, we discussed the impact of factors such as model size, training method, and architecture on LLMs' EQ. In summary, our study presents one of the first psychometric evaluations of the human-like characteristics of LLMs, which may shed light on the future development of LLMs aiming for both high intellectual and emotional intelligence. Project website: <https://emotional-intelligence.github.io/>

## Keywords

emotional intelligence, emotional understanding, Large Language Models, generative pretrained transformer, human-likeness, psychometric assessment

Received 28 July 2023; accepted 27 October 2023

## Introduction

Imagine an ancient male making a necklace from a pile of shells as a gift for a female. This endeavor would require at least two distinct types of abilities. First, he would need the foresight to conceptualize that if a hole were punched in each shell and a string threaded through these holes, the shells could form a necklace. Second, he must possess a rudimentary level of empathy, inferring that the female recipient of the necklace would likely experience joy. The former ability is a manifestation of the Systemizing Mechanism (Baron-Cohen, 2020), enabling humans to become the scientific and technological masters of our physical world. The latter, on the other hand, is referred to as Emotional Intelligence (EI), which allows us to think about our own and others' thoughts and feelings, thereby aiding us in navigating the social world (Mayer, Perkins et al., 2001; Mayer, Salovey et al., 2001; Salovey & Mayer, 1990).

In recent years, Large Language Models (LLMs) have made substantial strides, showcasing their expertise across multiple disciplines including mathematics, coding, visual comprehension, medicine, law, and psychology (Bubeck et al., 2023). Their impressive performance in logic-based tasks implies that LLMs, such as GPT-4, might be equipped with the Systemizing Mechanism comparable to human intelligence. Indeed, GPT-4 outperformed 99% of human participants in a modified text-based IQ test, a feat aligning

<sup>1</sup>Department of Psychology & Tsinghua Laboratory of Brain and Intelligence, Tsinghua University, Beijing, China

<sup>2</sup>Department of Psychology, Renmin University of China, Beijing, China

### Corresponding author:

Liu Jia, Department of Psychology & Tsinghua Laboratory of Brain and Intelligence, Tsinghua University, Beijing, China.

Email: [liujiaTHU@tsinghua.edu.cn](mailto:liujiaTHU@tsinghua.edu.cn)



with the elite MENSA (a high IQ society that provides a forum for intellectual exchanges among its members) level of general intelligence (King, 2023).

In contrast, investigations into the empathy of LLMs are relatively scarce and less systematic. Previous studies have mainly used the Theory of Mind (ToM) task, which measures the ability to understand and interpret others' mental states. LLMs launched before 2022 showed virtually no ability of ToM (Kosinski, 2023; Sap et al., 2023), whereas more recent models have shown significant improvement. For example, LLM "text-davinci-002" (January 2022) achieved an accuracy of 70%, comparable to that of six-year-old children, while LLM "text-davinci-003" (November 2022) reached 93%, on par with seven-year-old children (Kosinski, 2023). Specifically, the most advanced model, GPT-4, attained 100% accuracy with in-context learning (Moghaddam & Honey, 2023). While the ToM task provides valuable insights, it is not suitable to serve as a standardized test on EI for two reasons. First, ToM is a heterogeneous concept, spanning from false belief, the understanding that others can hold beliefs about the world that diverge from reality (Baron-Cohen et al., 1985), to pragmatic reasoning, the ability to incorporate contextual information and practical considerations when solving problems in real-world situations (Sperber & Wilson, 2002). Consequently, the heterogeneous nature of the ToM task may not meet the reliability and validity standards of psychometric tests. Second, the ToM task is simple for a typical human participant in general, rendering it more suitable to serve as a diagnostic tool for EI-related disorders such as the autism spectrum disorder rather than a discriminative test for general population. Consequently, standardized tests on EI, such as the Mayer-Salovey-Caruso Emotional Intelligence Test (MSCEIT, Mayer et al., 2003), do not include the ToM task.

According to EI theories (Mayer et al., 2016; Mayer & Salovey, 1995; Salovey & Mayer, 1990), emotion understanding (EU) is a fundamental component of EI, which serves as a subscale in MSCEIT. EU refers to the ability to recognize, interpret, and understand emotions in a social context, which lays the groundwork for effective communication, empathy, and social interaction (Mayer et al., 2016). Specifically, the test on EU is suitable for measuring the empathy of LLMs because they do not possess internal emotional states or experiences, and therefore they have to rely solely on accurately understanding and interpreting the social context to create more engaging and empathetic interactions.

In this study, we first developed a standardized EU test suitable for both humans and LLMs, termed the Situational Evaluation of Complex Emotional Understanding (SECEU). The SECEU is an objective, performance-driven, and text-based evaluation, providing a consistent assessment for both human and LLM

capabilities. The SECEU was designed in alignment with the EU subscale of MSCEIT (Mayer et al., 2003), a standardized test on EQ. To make the test suitable for LLMs, we removed the image recognition items in the MSCEIT, and greatly expanded the text-based items to capture a wider spectrum of emotional scenarios. Data from more than 500 young adults were collected to establish a norm for the SECEU. Then, with the SECEU, we evaluated a variety of mainstream and popular LLMs, including OpenAI GPT series (GPT-4, GPT-3.5-turbo, Curie, Babbage, DaVinci, text-davinci-001, text-davinci-002, and text-davinci-003), Claude, LLaMA-based models (Alpaca, Koala, LLaMA, and Vicuna), FastChat-T5, Pythia-based models (Dolly and Oasst), GLM-based models (ChatGLM), and RWKV (Recurrent Weighted Key-Value). Finally, we standardized the LLMs' scores against the norm, allowing for direct comparison with humans. We also compared the multivariate response patterns of the LLMs and human participants to compare their representation similarities.

## Methods

### Participants

A total of 541 human participants with valid responses were collected in this study. The participants ( $N = 541$ ; females: 339, males: 202; mean age: 22.33, SD: 2.49, ranging from 17 to 30 years) were all undergraduate and postgraduate college students in China. Informed consent was obtained prior to the SECEU test and participants were reimbursed after they completed the whole test. To ensure anonymity and data privacy, participants did not input any information that could identify them during the process. This study was approved by the Institutional Review Board at Tsinghua University.

We also invited three experts to take the SECEU test. Expert 1 is an accomplished Human Resources professional who has over 20 years of experience in navigating human emotions within diverse work environments, strengthening her discernment in EI. Expert 2 is a renowned figure in psychometrics and her expertise in creating tests assessing psychological variables lends exceptional rigor to our process. Expert 3 is an associate professor of psychology, whose deep understanding of human emotions, backed by her extensive academic achievements, makes her especially suitable for this test.

### Development of the SECEU

**Design.** Each item in the test presents a scenario grounded in the context of schools, families, or society. These scenarios, intricately designed with twists and turns, aim to evoke a mixture of positive and negative emotions. An illustrative example is: "Wang participated in a mathematics competition but felt he had not performed to his full

potential. However, when the results were announced, he found that he was in a position of top 10." These scenarios incorporate varying numbers of characters, and elicit emotions that can be self-directed, other-directed, or both. Participants are then presented with four of the most plausible emotions (e.g., surprised, joyful, puzzled, proud) associated with each scenario. They are then tasked to allocate a total of 10 points among these emotions based on the intensity they believe each emotion is experienced in the scenario. For instance, a response of 3, 3, 1, 3 suggests a multifaceted emotional response comprising 30% surprise, 30% joy, 10% puzzlement, and 30% pride. The test contains a total of 40 items.

**Scoring and standardization.** Building on the premise that groups possess collective emotional knowledge (Legree et al., 2005), we utilized a consensus scoring method (Palmer et al., 2005) to standardize the SECEU test. We calculated a standard score for each emotion in each item by averaging the corresponding scores from human participants, resulting in a  $40 \times 4$  symmetric matrix (see <https://emotional-intelligence.github.io/> for the standard score of humans, SS). Each participant's EU ability was measured by calculating the Euclidean distance between their individual scores and the group-derived standard scores for each item. Smaller distances signify higher EU abilities. Reliability was affirmed by assessing the internal consistency (Cronbach's  $\alpha$ ) of participants' performance across all 40 items, and validity was confirmed by EI experts' evaluations.

## Procedure

**The SECEU test for human participants.** The online SECEU test was built on the JATOS platform (Lange et al., 2015) based on the Jspsych plugin (de Leeuw et al., 2023), which was written in the React Framework (<https://reactjs.org/>). Each item was presented to the participants with a scenario and followed by four emotion options (40 items in total, see <https://emotional-intelligence.github.io/> for both English and Chinese versions). Participants were instructed to read the scenario and then allocate a total of 10 points across the four emotion options based on the intensity of each emotion experienced by the person in the scenario. There were no correct or incorrect answers. Participants were encouraged to respond according to their own understanding and interpretation of the scenarios.

**The SECEU test for LLMs.** A variety of mainstream LLMs, including the OpenAI GPT series (GPT-4, GPT-3.5-turbo, Curie, Babbage, DaVinci, text-davinci-001, text-davinci-002, and text-davinci-003), Claude, LLaMA-based models (Alpaca, Koala, LLaMA, and Vicuna), FastChat-T5, Pythia-based models (Dolly and Oasst), GLM-based models (ChatGLM), and RNN-based

models (RWKV), were evaluated by the SECEU test. Given that the majority of these models are primarily trained on English datasets, using the English version of the SECEU provides a more accurate assessment of their performance, allowing for a clearer comparison between their abilities and those of a human. As a result, the English version of SECEU was presented to the LLMs instead of the Chinese version.

The task was in a direct question-and-answer format. We asked, for example, "Story: Wang participated in a mathematics competition but felt he had not performed to his full potential. However, when the results were announced, he found that he was in a position of top 10. He would feel: Options: (1) Surprised; (2) Joyful; (3) Puzzled; (4) Proud. Assign a score to each option based on the story, sum up to 10." There could be very subtle changes of the direct prompt. For instance, we used "provide a score for each emotion based on the emotion (sum of four options should be of 10 points)" for Dolly. There were a set of paraphrases of the direct prompt to get the best performance.

To decrease the randomness, a constant temperature parameter was set to 0.1 and the top\_p parameter was set to 1 across all these models. A lower temperature was set to make LLMs' output more focused and deterministic to achieve a higher reliability in responding to the test. To dictate the maximum length of the generated text, the max\_tokens parameter was set to 512.

Before being processed by the models, text data underwent several preprocessing steps to normalize it. This normalization process ensures that data fed into the models are in a consistent and appropriate format, enhancing the output's quality. If a model did not provide any meaningful response to an item, the response for this item was predefined as a null vector (0, 0, 0, 0). A few models failed to generate a response for a majority of items (LLaMA: 40; FastChat-T5: 31; RWKV-v4: 31; DaVinci: 40; Curie: 40; Babbage: 40; text-davinci-001: 26; text-davinci-002: 28; marked as "FAILED" in Table S1). Additionally, several models were unable to provide the answer which the summation of the four scores equaled 10:

- (i) Null Responses: Answer vectors signifying null responses, i.e., (0, 0, 0, 0), were preserved as such (Alpaca: 1; ChatGLM: 1).
- (ii) Handling Negative Values: For datasets encompassing negative values, an addition operation involving the absolute value of the lowest number was performed across all elements, followed by a subsequent normalization to maintain consistency with the original scale. For instance, an original data vector of (-4, -2, -2, 2) would be adjusted to (0, 2, 2, 6).
- (iii) Normalization of Answer Vectors: The remaining answer vectors were normalized to achieve a cumulative score of 10. This involved proportionally distributing a score of 10 among the answer vector based

on the contribution of each value to the total score on this item.

**LLMs' EQ.** The performance of each LLM was standardized by calculating the Euclidean distance between the

model's responses (LLM) and the standard scores of humans (SS) on each item  $i$  (from 1 to 40) and then averaged to yield the SECEU score. Lower SECEU scores indicate greater alignment with human standards.

$$\text{SECEU score} = \frac{1}{40} \sum_{i=1}^{40} \sqrt{(\text{LLM}_{i1} - \text{SS}_{i1})^2 + (\text{LLM}_{i2} - \text{SS}_{i2})^2 + (\text{LLM}_{i3} - \text{SS}_{i3})^2 + (\text{LLM}_{i4} - \text{SS}_{i4})^2} \quad (1)$$

The SECEU score was then normalized into an EQ score which was designed to follow a normal distribution with the average score set at 100 and the standard deviation at 15. The standardization process involved the following steps: (1) the original SECEU score was subtracted from the mean value of the SECEU scores of the human norm ( $M$  in equation (2)) and divided by the standard deviation of the SECEU scores of the human norm ( $SD$  in equation (2)), and (2) the resulting value was then multiplied by the new standard deviation (15) and added to the new mean value (100), yielding the EQ score. Thus, the EQ score represents a normalized measure of the LLM's EQ, permitting easier comparison across different models.

$$\text{LLM's EQ} = 15 \times \frac{M - \text{SECEU score}}{SD} + 100 \quad (2)$$

**LLMs' representational pattern.** To establish strong equivalence between the LLMs and humans, we examined whether they employed similar representations to conduct the test. Item-wise correlation analysis (Izard & Spelke, 2009; Tian et al., 2020) was applied to compare response patterns between the LLMs and human participants. The human template (a vector with a length of 40; see <https://emotional-intelligence.github.io/> for the human pattern template) was generated by averaging the multi-item discriminability patterns across all human participants, where each pattern was constructed based on the distance of each item to the standard score. The multi-item discriminability pattern of a specific LLM was also calculated based on the distance of each item  $i$  (from 1 to 40) to the standard scores of humans (SS).

$$\text{Discriminability}_i = \sqrt{(\text{LLM}_{i1} - \text{SS}_{i1})^2 + (\text{LLM}_{i2} - \text{SS}_{i2})^2 + (\text{LLM}_{i3} - \text{SS}_{i3})^2 + (\text{LLM}_{i4} - \text{SS}_{i4})^2} \quad (3)$$

We calculated the Pearson correlation coefficient between the discriminability pattern of each participant and the human template (Human-to-Human similarity). To avoid the inflation in calculating correlation, the template was constructed excluding the individual whose Human-to-Human correlation coefficient was calculated. Specifically, for participant 001, the Human-to-human correlation coefficient was calculated

as the Pearson correlation between his/her own discriminability pattern and the average discriminability pattern of the remaining 540 participants. The distribution of the Human-to-Human similarity served as a norm for pattern similarity. The Pearson correlation coefficient between the discriminability pattern of each LLM and the human template was calculated as the LLM-to-Human similarity.

$$\text{LLM - to - Human similarity}_{\text{Pearson}} = \frac{\sum_{i=1}^{40} (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^{40} (X_i - \bar{X})^2} \sqrt{\sum_{i=1}^{40} (Y_i - \bar{Y})^2}} \quad (4)$$

Here,  $X_i$  and  $Y_i$  represent the item  $i$  of the "Discriminability" vector and the human template vector, respectively. The length of both vectors is 40.  $\bar{X}$  and  $\bar{Y}$  denote the mean of  $X_i$  and  $Y_i$ , respectively. If the LLM-to-Human similarity is less than one SD of Human-to-human similarity below the population, such LLM is considered as employing a qualitatively different mechanism from humans in EU.

**Prompt engineering.** Prompt engineering—the meticulous development and choice of prompts—plays a pivotal role in the efficacy of LLMs (Hebenstreit et al., 2023; Hendrycks et al., 2021; Nair et al., 2023; OpenAI, 2023; Shinn et al., 2023). In essence, prompt engineering refers to the strategy of designing and selecting prompts that can substantially guide and influence the responses of LLMs. The necessity of prompt engineering lies in its potential to enhance the precision and relevance of the responses generated by these models, thereby leading to more effective and reliable outcomes. In the realm of EI, prompts serve a crucial function. They provide a direction for the model, enabling it to understand and generate responses that are not just accurate but also emotionally intelligent. Given the nature of EI that involves understanding, processing, and managing emotions, prompts can significantly aid the LLMs in identifying the correct context and producing responses that exhibit emotional understanding and empathy.

To examine the prompts' influence on EU ability, thereby optimizing model outputs, four kinds of prompt engineering techniques (see <https://emotional-intelligence.github.io/> for prompts) were applied to the OpenAI GPT series (GPT-3.5-turbo, Curie, Babbage, DaVinci, text-davinci-001, text-davinci-002, and text-davinci-003): (1) Zero-shot Prompts, (2) Enhanced Zero-shot Prompts Incorporating Step-by-Step Thinking, (3) Two-shot Chain of Thought Reasoning Approaches, and (4) Two-shot Chain of Thought Reasoning Approaches Augmented with Step-by-Step Thinking.

To decrease the randomness, a constant temperature parameter was set to 0 and the top\_p parameter was set to 0.9 across all these models. To dictate the maximum length of the generated text, the max\_tokens parameter was set to 2048. The normalization process was the same as the one without prompts.

### Data availability

The SECEU test (both English and Chinese Versions), the code for the test on human participants, the standardized scores, the norm, and the prompts are available at <https://emotional-intelligence.github.io/>. The raw data of human participants are available from the corresponding author upon request.

## Results

### The development of a standardized test on EU

The SECEU is designated to measure EU, which comprises 40 items (see <https://emotional-intelligence.github.io/> for both English and Chinese versions). Each item describes a scenario set either in a school, family, or social contexts with twists and turns designed to evoke a mixture of positive and negative emotions (e.g., "Wang participated in a mathematics competition but felt he had not performed to his full potential. However, when the results were announced, he found that he was in a position of top 10"). The scenarios feature a varying number of characters, and emotions can be self-directed, other-directed, or both. For each scenario, four of the most plausible emotions (e.g., surprised, joyful, puzzled, proud) are listed. Participants were asked to evaluate the intensity of each emotion with numbers that were added up to 10 (e.g., 3, 3, 1, 3, indicating a multifaceted emotion response comprising 30% surprise, 30% joy, 10% puzzlement, and 30% pride). Figure 1 shows two items of the SECEU test and the standard scores by averaging answers across the participants.

Under the assumption that groups possess collective knowledge of emotion (Legree et al., 2005), we adopted a consensus scoring method to standardize the SECEU (Palmer et al., 2005). To do this, we administered the SECEU to a large sample of undergraduate and postgraduate students ( $N = 541$ ; females: 339, males: 202; mean age: 22.33, SD: 2.49, ranging from 17 to 30 years). Then, we calculated the standard score for each emotion of each item by averaging corresponding scores across the participants. Accordingly, we measured each participant's EU ability on each item by calculating the Euclidean distance between the participant's individual score and the standard score derived from the whole group for each item, with smaller distances indicating better EU ability. This analysis revealed significant variance in individual differences in EU ability ( $M = 2.79$ ,  $SD = 0.822$ , range from 1.40 to 6.29), suggesting that the SECEU is well suited to serve as a discriminative test for assessing EU in a general population.

To evaluate the reliability of the SECEU, we assessed the internal consistency of participants' performance (i.e., the Euclidean distance) across the 40 items, and revealed a high reliability of the test (Cronbach's  $\alpha = 0.94$ ). We further examined the distribution of participants' performance on each item (Figure S1) and found no evidence of ceiling or floor effects, with mean distances varying from 2.19 to 3.32 and SD ranging from 1.13 to 1.82. In addition, there was no significant sex difference (male: 2.85, female: 2.76,  $t(539) = 1.34$ ,  $p = 0.18$ ).

To evaluate the validity of the SECEU, we invited three experts known for their high EI to take the test. All experts' performance exceeded at least 73% of the population of all

The SECEU test		The Standard Scores:			
<b>Item 1</b>	<b>Options:</b>	(1) Expectant	(2) Excited	(3) Joyful	(4) Frustrated
Story: The airplane model that Wang made fell from the sky one minute after take-off. When she inspected the model, she found a part that could possibly be improved. At this moment, she would feel:		3.56	3.09	1.78	1.57
<b>Item 2</b>	<b>Options:</b>	(1) Desperate	(2) Fear	(3) Helpless	(4) Sad
Story: Although Aunt Li pays close attention to her lifestyle and diet, she still got a malignant tumor. The chances of curing the tumor now are quite slim, and Aunt Li can do nothing about it. She would feel:		3.45	2.04	2.17	2.33
⋮				⋮	
		40 × 4 matrix			

**Figure 1.** SECEU test items and the standard scores from the population. For the whole set of the test, see <https://emotional-intelligence.github.io/>.

participants in this study, indicating that the test is effective in differentiating experts from the general population. Specifically, the average score of the three experts exceeded 99% of the whole population tested, further confirming the validity of using the consensus scoring method in standardizing the SECEU.

Finally, we constructed the norm for EU by converting participants' raw scores in the SECEU into standard EQ (Emotional Quotient) scores, designed to follow a normal distribution with the average score set at 100 and standard deviation at 15. In practical terms, an individual with an EQ of 100 possesses an EU ability corresponding to the population average. Meanwhile, an individual with an EQ of 115 outperforms approximately 84% of the population (i.e., one SD above the population average), and an individual with an EQ score of 130 exceeds 97.7% of the population.

### The assessment of LLMs' EQ

We evaluated a variety of mainstream LLMs using the SECEU, and then standardized their scores based on the norm of the human participants for a direct comparison between LLMs and humans. These LLMs included OpenAI GPT series (GPT-4, GPT-3.5-turbo, Curie, Babbage, DaVinci, text-davinci-001, text-davinci-002,

and text-davinci-003), Claude, as well as open-source models such as LLaMA-based models (Alpaca, Koala, LLaMA, and Vicuna), Pythia-based models (Dolly and Oasst), GLM-based models (ChatGLM), and FastChat-T5. Recurrent Weighted Key-Value (RWKV), which utilizes recurrent neural networks (RNNs) instead of transformers, was also included.

Some models, including LLaMA, FastChat-T5, and RWKV-v4, were unable to complete the test even with the assistance of prompts (Table 1). A few LLMs, including DaVinci, Curie, Babbage, text-davinci-001, and text-davinci-002 managed to complete the test with prompts such as Two-shot Chain of Thought (COT) and Step-by-Step prompts (see Supplementary for the prompt engineering). In addition, other models, such as text-davinci-003, were able to complete the test but performance was significantly improved with prompts. Here, we only included models' best performance to examine how closely they can approach human-level performance under ideal conditions (Table 1; see also Tables S1 and S2). To directly compare to human participants, the performance of each model was standardized by calculating the Euclidean distance between the model's responses and the standard scores of humans, which was then normalized into an EQ score (Table 1). Finally, these

**Table 1.** LLMs' EQ, representational patterns, and properties

Based	Models	SECEU score	EQ		Pattern Similarity		Properties			
			EQ	%	<i>r</i>	%	Size	Time	SFT	RLHF
OpenAI	GPT series									
	DaVinci #	3.5	87	18%	0.41**	91%	175B	2020/05	×	×
	Curie #	2.7	102	50%	0.11	29%	13B	Unknown	×	×
	Babbage #	2.78	100	44%	−0.12	4%	3B	Unknown	×	×
	text-davinci-001 #	2.4	107	64%	0.2	47%	<175B	Unknown	×	×
	text-davinci-002 #	3.3	91	23%	−0.04	8%	<175B	Unknown	✓	×
	text-davinci-003 ###	2.01	114	83%	0.31*	73%	175B	2022/11/28	✓	✓
	GPT-3.5-turbo	2.63	103	52%	0.04	17%	175B	2022/11/30	✓	✓
	GPT-4	1.89	117	89%	0.28	67%	Unknown	2023/03/14	✓	✓
LLaMA	LLaMA		—FAILED—				13B	2023/02/24	×	×
	Alpaca	2.56	104	56%	0.03	15%	13B	2023/03/09	✓	×
	Vicuna	2.5	105	59%	−0.02	10%	13B	2023/03/30	✓	×
	Koala	3.72	83	13%	0.43**	93%	13B	2023/04/03	✓	×
Flan-T5	FastChat-T5		—FAILED—				3B	2023/04/30	✓	×
Pythia	Dolly	2.89	98	38%	0.26	62%	13B	2023/04/12	✓	×
	Oasst	2.41	107	64%	0.24	59%	13B	2023/04/15	✓	✓
GLM	ChatGLM	3.12	94	28%	0.09	24%	6B	2023/03/14	✓	✓
RWKV	RWKV-v4		—FAILED—				13B	2023/02/15	✓	×
Claude	Claude	2.46	106	61%	0.11	28%	Unknown	2023/03/14	✓	✓

Notes: Table shows the SECEU scores, EQ scores, representational pattern similarity, and properties of mainstream LLMs evaluated in the current study. #: models require prompts to complete the test; ##: models' performance benefits from prompts; FAILED: even with prompts, the LLMs cannot complete the test; %: the percent of humans whose performance was below that of an LLM in the test; Pattern Similarity: the degree of similarity is indexed Pearson correlation coefficient (*r*); \*:  $p < 0.05$ ; \*\*:  $p < 0.01$ ; Size: the parameter size of LLMs in the unit of billions (B); Time: the launch time in the format YYYY/MM/DD; SFT: supervised fine-tune; RLHF: reinforcement learning from human feedback; ✓: yes; ×: no.

LLMs were categorized as either expert (above 115), normal (between 85 and 115), or poor (below 85) based on their EQ scores.

The results revealed a substantial variation in EU among the LLMs tested (Figure 2). Within the OpenAI GPT series, GPT-4 achieved the highest EQ of 117, exceeding 89% of humans. In contrast, DaVinci scored the lowest, with an EQ of 87, only outperforming 18% of humans.

The LLaMA-based models generally scored lower than the OpenAI GPT series, with Alpaca and Vicuna achieving the highest EQ of 104 and 105, respectively. Conversely, Koala showed the poorest performance, with an EQ score of 83, only surpassing 13% of humans. The base model LLaMA was unable to complete the test. Other models, such as Oasst (EQ: 107), Dolly (EQ: 98), ChatGLM (EQ: 94), and Claude (EQ: 106), all fell within the normal range.

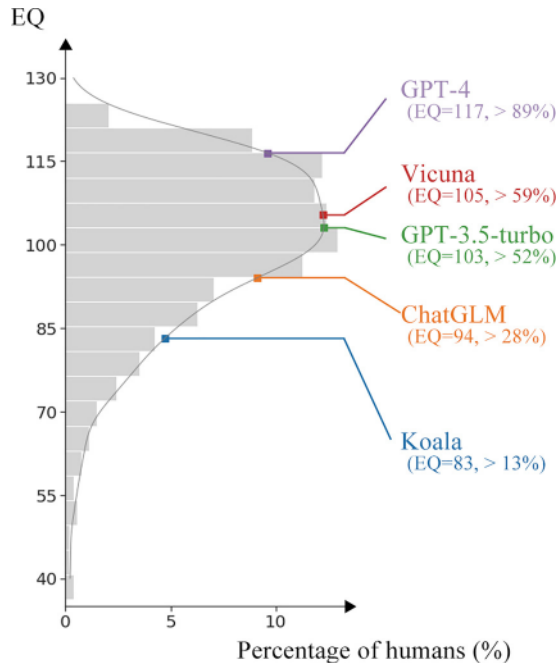
In short, the majority of the LLMs tested showed satisfactory EU scores, comparable to those of the average human population. Specifically, GPT-4 reached the expert level of humans.

### The assessment of LLMs' representational pattern

The measurement of the LLMs' EQ scores provides an index of their EU ability within the reference frame of humans. A further question is whether they employ human-like mechanisms to evaluate complex emotions in scenarios. The univariate analysis used to compare EQ scores between human participants and LLMs only suggests weak equivalence, as a model could achieve a high EQ score using mechanisms that qualitatively differ from humans. Therefore, to establish strong equivalence between the LLMs and humans, we examined whether they employed similar representations to conduct the test.

One approach is to use the item-wise correlation analysis (Izard & Spelke, 2009; Tian et al., 2020) to compare response patterns between the LLMs and human participants. To do this, we first constructed a multi-item discriminability vector (i.e., an item-wise response pattern) for each participant by using the distance of each item to the standard score, and thus this vector's length corresponded to the

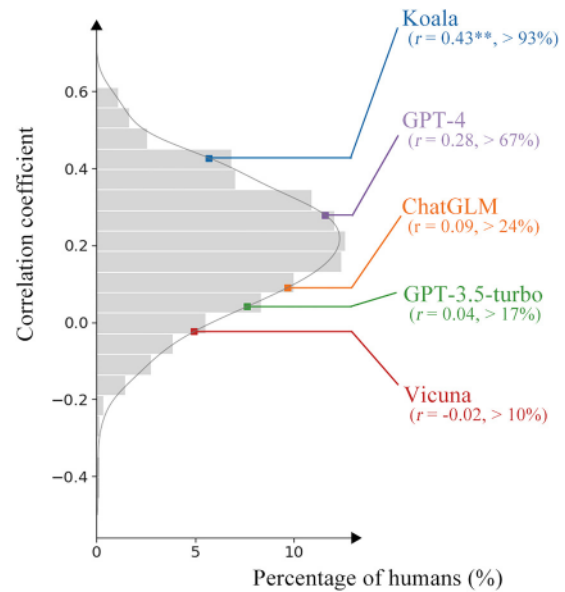




**Figure 2.** LLMs' EQ. The light-grey histogram represents the distribution of human participants' EQ scores, with the y-axis indicating the EQ score and the x-axis showing the percentage of total participants. The grey kernel density estimation (KDE) line demonstrates the probability density of the EQ scores. Key points are highlighted with colored square markers for LLMs (e.g., GPT-4's EQ score is 117, marked by the purple square, exceeding 89% of the human participants). For simplicity, here we only present the performance from GPT-4, Vicuna, GPT-3.5-turbo, ChatGLM, and Koala.

number of items (i.e., 40). Then, we created a template of response patterns by averaging the multi-item discriminability patterns across all human participants, along with the distribution of the correlation coefficients between each participant's response pattern and the template (Human-to-Human similarity) to serve as a norm for pattern similarity ( $M=0.199$ ,  $SD=0.166$ ). Finally, we quantified the similarity in response patterns between the LLMs and humans by calculating the correlation coefficient between the multi-item discriminability vector of each LLM and the human template (LLM-to-Human, Table 1). An LLM that has an LLM-to-Human correlation coefficient one SD deviation below the mean of Human-to-Human distribution is considered as employing a qualitatively different mechanism from humans in EU.

Surprisingly, despite its lower performance in the SECEU, Koala showed the highest similarity in representational patterns to humans ( $r=0.43$ ,  $p<0.01$ , exceeding 93% of human participants) (Figure 3). This suggests that Koala may represent emotions in the same way as humans do, as it performed well on items where humans excelled and struggled on items where humans faced

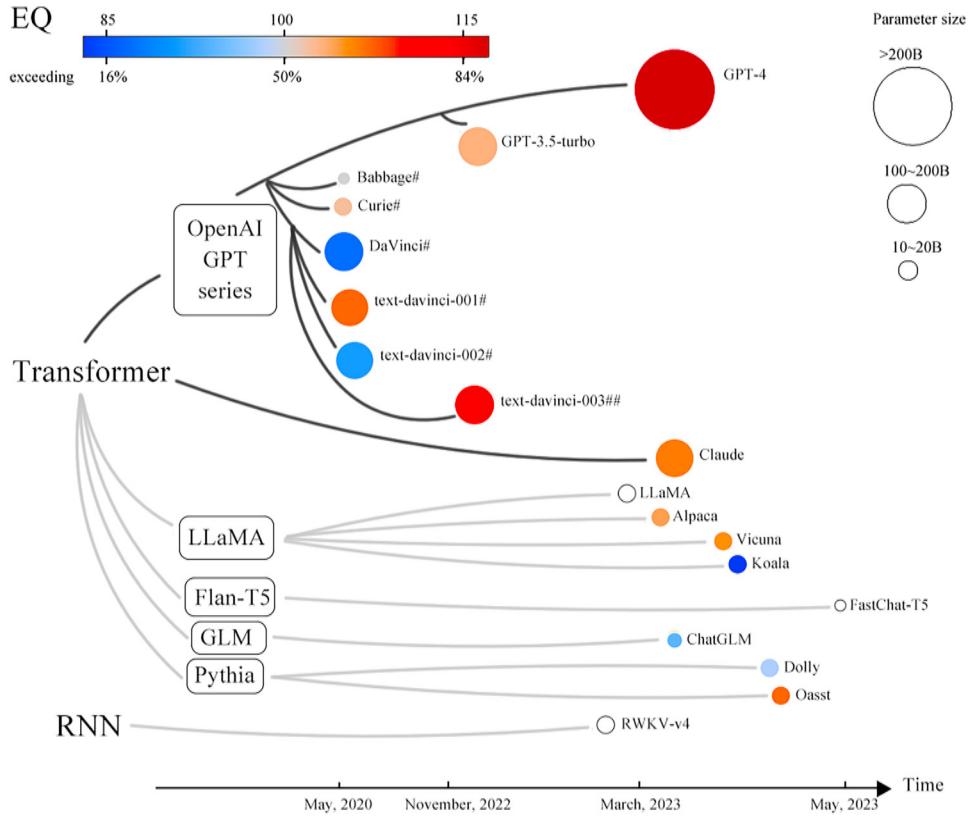


**Figure 3.** The pattern similarity between LLMs and humans. The light-grey histogram represents the distribution of Human-to-Human pattern similarity, with the y-axis indicating the Pearson correlation coefficients and the x-axis showing the percentage of total participants. The KDE line demonstrates the probability density of the Pearson correlation coefficient. Key points are highlighted with colored square markers for LLMs. For simplicity, here we only present the performance from GPT-4, Vicuna, GPT-3.5-turbo, ChatGLM, and Koala. \*\*:  $p<0.01$ .

challenges. That is, the discrepancies in understanding emotions between Koala and humans are rather quantitative than qualitative. On the other hand, the representational patterns of models such as Babbage, text-davinci-002, Alpaca, and Vicuna differed qualitatively from humans' representational patterns (Babbage:  $r=-0.12$ , >4%; text-davinci-002:  $r=-0.04$ , >8%; Alpaca:  $r=0.03$ , >15%; Vicuna:  $r=-0.02$ , >10%). This suggests that, despite their above-average EQ scores, these models likely employed mechanisms that are qualitatively different from human processes.

GPT-4, the most advanced model to date, showed high similarity in representational pattern ( $r=0.28$ , >67%). This result implies that GPT-4 may have significantly changed its architecture or implemented novel training methods to align its EU ability more closely to humans. Interestingly, prompts apparently played a critical role in improving representational similarity. With two-shot COT prompts, DaVinci and text-davinci-003 showed high similarity in representational pattern to humans (Davinci:  $r=0.41$ ,  $p<0.01$ , >91%; text-davinci-003:  $r=0.31$ ,  $p<0.05$ , >73%), higher than that of GPT-4. Note that without prompts, they failed to complete the SECEU test. In contrast, prompts had little effect on GPT-4 and GPT-3.5-turbo.





**Figure 4.** The family tree of LLMs. Each node in the tree represents an LLM, whose vertical position along the x-axis indicates the launch time. The size of each node corresponds to the parameter size of the LLM. Note that the size of GPT-4 and Claude was estimated based on publicly available information. Color denotes the EQ scores, with red color for higher scores and blue color for lower scores. Note that the uncolored circles show that models failed to conduct the SECEU. The color of the branches distinguishes between open-source (light gray) and closed-source (dark gray) models.

## Discussion

Since the debut of ChatGPT, a great number of tasks and benchmarks have been developed to examine the capacities. These empirical evaluations and analyses mainly focus on language generation (e.g., conditional text generation), knowledge utilization (e.g., closed-book and open-book QAs), and complex reasoning (e.g., symbolic and mathematical reasoning) (Zhao et al., 2023). However, tests on human alignment of LLMs to human values and needs, a core ability for the broad use of LLMs in the real world, are relatively scarce. Here in this study, we used traditional psychometric methods to develop a valid and reliable test on emotional understanding, the SECEU, to evaluate the EI of LLMs. We found that a majority of the LLMs tested performed satisfactorily in the test, achieving above-average EQ scores, although significant individual differences were present across the LLMs. Critically, some LLMs apparently did not reply on the human-like representation to achieve human-level performance, as their representational patterns diverged significantly from human patterns, suggesting a qualitative difference in the underlying mechanisms. In summary, our study provides the

first comprehensive psychometric examination of the EI of LLMs, which may shed light on the development of future LLMs that embody high levels of both intellectual and emotional intelligence.

Various factors appear to influence the EQ scores of the LLMs (Figure 4). The most conspicuous one is the model size, which is essential to emergent abilities (Bubeck et al., 2023), making AI algorithms unprecedentedly powerful and effective. While the larger models generally scored higher in the test, certain smaller models such as Oasst and Alpaca still managed to achieve satisfactory EQ scores. This suggests that factors beyond the mere size may have a more profound influence on models' EU.

The effectiveness of various training methods, such as supervised training, reinforcement learning, self-supervised learning, and a combination thereof, likely substantially influences the EQ scores. For example, despite architectural differences (Pythia versus LLaMA), Oasst and Alpaca yielded similar scores in the test, demonstrating the potential of well-implemented fine-tuning techniques. In fact, these enhancements may be achieved through two main avenues. The first involves supervised fine-tuning (SFT),

which allows for more structured and targeted fine-tuning, thereby improving models' linguistic ability and their grasp of contextual nuances (Köpf et al., 2023; Taori et al., 2023a). The other approach employs reinforcement learning from human feedback (RLHF), enabling the models to learn from human insights and thereby fostering more human-like responses. Indeed, there is a giant leap in EU seen between text-davinci-002 (>23%) to text-davinci-003 (>83%), two different versions of the same model with the latter employing RLHF.

Another influential factor is the model architecture. Models using the Transformer architecture, such as the GPT series and the LLaMA-based models, generally performed well in this test. In contrast, models using RNNs, such as RWKV-v4, failed to complete the test even with the help of various prompts. Besides, within the Transformer architecture, the "decoder-only" or causal decoder models (e.g., the GPT series), which generate sequences based solely on a self-attention mechanism (Brown et al., 2020; Vaswani et al., 2017), outperformed the "encoder-decoder" models (e.g., FastChat-T5), which incorporate an extra step to interpret input data into meaningful representations (Devlin et al., 2019; Zheng et al., 2023).

In summary, our study provides novel evaluation on the human-like characteristics of LLMs, along with the tests on self-awareness (Kosinski, 2023) and affective computing (Amin et al., 2023). However, because only a limited number of LLMs were tested in this study (results on more LLMs will be continuously updated in <https://emotional-intelligence.github.io/>), our findings are biased and inconclusive. Further, there are more questions that need to be explored in future studies. First, this study focused solely on the EU ability of the LLMs, while EI is a multifaceted construct encompassing not only EU but also emotion perception, emotion facilitation, and emotion management (e.g., Mayer et al., 2016; Mayer & Salovey, 1995; Salovey & Mayer, 1990). Therefore, future studies could design scenarios to examine whether LLMs can assist humans in leveraging emotions to facilitate cognitive processes and effectively manage their emotional responses.

Second, EI requires the integration of various facets to execute complex tasks, which necessitate not only an understanding of emotions, but also the comprehension of thoughts, beliefs, and intentions. Future studies should adopt broader scope assessments, akin to ToM tests, while avoiding their lack of discriminative power. Besides, with recent advancements, LLMs are now capable of processing multimodal information (Gao et al., 2023; Wang et al., 2023; Yang et al., 2023; Zhu et al., 2023). Therefore, future studies should investigate how LLMs interpret complex emotions from multimodal inputs, such as text combined with facial expressions or the tone of voice. In short, tests that combine emotions

with cognitive factors based on multimodal clues likely furnish a more comprehensive understanding of LLMs' EI, which is critical for LLMs' effective and ethically responsible deployment in real-world scenarios of mental health, interpersonal dynamics, work collaboration, and career achievement (e.g., Dulewicz & Higgs, 2000; Hanafi & Noor, 2016; Lea et al., 2019; Mayer et al., 2016; McCleskey, 2014; Warwick & Nettelbeck, 2004).

Third, although our study introduces a novel perspective on the comparative behaviors of LLMs and humans in emotional understanding, the test containing only 40 items provides a limited view and does not fully encapsulate comprehensive characteristics of LLMs and human emotional understanding. Future studies could benefit from larger and more diverse tests, along with multifaceted analytical approaches, to delve deeper into qualitative differences in emotional understanding between LLMs and humans.

Finally, while the factors examined in this study contribute to our standing of LLM's EU, they are largely descriptive and thus do not establish causal relationships. With the recent progress of open-source LLMs (Bai et al., 2022; Chiang et al., 2023; Conover et al., 2023; Geng et al., 2023; Köpf et al., 2023; Peng et al., 2023; Taori et al., 2023b; Touvron et al., 2023; Zeng et al., 2022; Zheng et al., 2023), direct manipulation of the potentially influential factors, such as training approaches and model size, has become plausible. Such manipulations will facilitate the establishment of causal relationships between these factors and models' EI ability, offering valuable insights for the development of future LLMs with better EI.

### Acknowledgements

This study was funded by Shuimu Tsinghua Scholar Program (X.W.), Beijing Municipal Science & Technology Commission, Administrative Commission of Zhongguancun Science Park (Z221100002722012), and Tsinghua University Guoqiang Institute (2020GQG1016).

### Authors' contribution

J.L. conceived the study and provided advice. X.L. developed the SECEU test, and X.L. and X.W. translated it into English. Z.Y. built the online SECEU test, and X.W. carried out the SECEU test for human participants, built the norm of EQ, and analyzed the data. Y.W. performed the SECEU test for LLMs, and Z.Y. wrote the prompts. X.W. and J.L. wrote the manuscript with suggestions and revisions from X.L., Z.Y., and Y.W.

### Declaration of conflicting interests


The authors declared no potential conflicts of interest with respect to the research, authorship, and/or publication of this article.

### Funding

The authors disclosed receipt of the following financial support for the research, authorship, and/or publication of this article: This

work was supported by the Tsinghua University Guoqiang Institute, Shuimu Tsinghua Scholar Program, Beijing Municipal Science & Technology Commission, Administrative Commission of Zhongguancun Science Park (grant number 2020GQG1016, Z221100002722012).

## ORCID iD

Xuena Wang  <https://orcid.org/0000-0003-3703-2674>

## Supplemental material

Supplemental material for this article is available online.

## References

- Amin, M. M., Cambria, E., & Schuller, B. W. (2023). *Will affective computing emerge from foundation models and general AI? A first evaluation on ChatGPT*. <https://doi.org/10.48550/ARXIV.2303.03186>
- Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., Chen, A., Goldie, A., Mirhoseini, A., McKinnon, C., Chen, C., Olsson, C., Olah, C., Hernandez, D., Drain, D., Ganguli, D., Li, D., Tran-Johnson, E., & Perez, E., ... J. Kaplan (2022). *Constitutional AI: Harmlessness from AI feedback* (arXiv:2212.08073). arXiv. <http://arxiv.org/abs/2212.08073>
- Baron-Cohen, S. (2020). *The pattern seekers: How autism drives human invention*. Basic Books.
- Baron-Cohen, S., Leslie, A. M., & Frith, U. (1985). Does the autistic child have a "theory of mind"? *Cognition*, 21(1), 37–46. [https://doi.org/10.1016/0010-0277\(85\)90022-8](https://doi.org/10.1016/0010-0277(85)90022-8)
- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., Agarwal, S., Herbert-Voss, A., Krueger, G., Henighan, T., Child, R., Ramesh, A., Ziegler, D., Wu, J., Winter, C., ... & Amodei, D. (2020). Language models are few-shot learners. Proceedings of the 34th International Conference on Neural Information Processing Systems, Vancouver, BC, Canada.
- Bubeck, S., Chandrasekaran, V., Eldan, R., Gehrke, J., Horvitz, E., Kamar, E., Lee, P., Lee, Y. T., Li, Y., Lundberg, S., Nori, H., Palangi, H., Ribeiro, M. T., & Zhang, Y. (2023). *Sparks of artificial general intelligence: Early experiments with GPT-4*. <https://doi.org/10.48550/ARXIV.2303.12712>
- Chiang, W.-L., Li, Z., Lin, Z., Sheng, Y., Wu, Z., Zhang, H., Zheng, L., Zhuang, S., Zhuang, Y., Gonzalez, J. E., Stoica, I., & Xing, E. P. (2023). *Vicuna: An open-source Chatbot impressing GPT-4 with 90%\* ChatGPT quality*. <https://lmsys.org/blog/2023-03-30-vicuna/>
- Conover, M., Hayes, M., Mathur, A., Meng, X., Xie, J., Wan, J., Shah, S., Ghodsi, A., Wendell, P., Zaharia, M., & Xin, R. (2023). *Free dolly: Introducing the world's first truly open instruction-tuned LLM*. URL <https://www.databricks.com/blog/2023/04/12/dolly-first-open-commercially-viable-instruction-tuned-llm>
- de Leeuw, J. R., Gilbert, R. A., & Luchterhandt, B. (2023). Jpsy: Enabling an open-source collaborative ecosystem of behavioral experiments. *Journal of Open Source Software*, 8(85), 5351. <https://doi.org/10.21105/joss.05351>
- Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019). *BERT: Pre-training of deep bidirectional transformers for language understanding* (arXiv:1810.04805). arXiv. <http://arxiv.org/abs/1810.04805>
- Dulewicz, V., & Higgins, M. (2000). Emotional intelligence—A review and evaluation study. *Journal of Managerial Psychology*, 15(4), 341–372. <https://doi.org/10.1108/02683940010330993>
- Gao, P., Han, J., Zhang, R., Lin, Z., Geng, S., Zhou, A., Zhang, W., Lu, P., He, C., Yue, X., Li, H., & Qiao, Y. (2023). *LLaMA-adaptor: Parameter-efficient visual instruction model*. <https://doi.org/10.48550/ARXIV.2304.15010>
- Geng, X., Gudibande, A., Liu, H., Wallace, E., Abbeel, P., Levine, S., & Song, D. (2023). *Koala: A dialogue model for academic research*. <https://bair.berkeley.edu/blog/2023/04/03/koala/>
- Hanafi, Z., & Noor, F. (2016). Relationship between emotional intelligence and academic achievement in emerging adults: A systematic review. *International Journal of Academic Research in Business and Social Sciences*, 6(6), 268–290. <https://doi.org/10.6007/IJARBS/v6-i6/2197>
- Hebenstreit, K., Praas, R., Kiesewetter, L. P., & Samwald, M. (2023). *An automatically discovered chain-of-thought prompt generalizes to novel models and datasets* (arXiv:2305.02897). arXiv. <http://arxiv.org/abs/2305.02897>
- Hendrycks, D., Burns, C., Basart, S., Zou, A., Mazeika, M., Song, D., & Steinhardt, J. (2021). *Measuring massive multitask language understanding* (arXiv:2009.03300). arXiv. <http://arxiv.org/abs/2009.03300>
- Izard, V., & Spelke, E. S. (2009). Development of sensitivity to geometry in visual forms. *Human Evolution*, 23(3), 213–248.
- King, M. (2023). *Administration of the text-based portions of a general IQ test to five different large language models* [Preprint]. <https://doi.org/10.36227/techrxiv.22645561.v1>
- Köpf, A., Kilcher, Y., von Rütte, D., Anagnostidis, S., Tam, Z.-R., Stevens, K., Barhoum, A., Duc, N. M., Stanley, O., Nagyi, R., ES, S., Suri, S., Glushkov, D., Dantuluri, A., Maguire, A., Schuhmann, C., Nguyen, H., & Mattick, A. (2023). *OpenAssistant conversations—Democratizing large language model alignment*. ArXiv Preprint ArXiv:2304.07327.
- Kosinski, M. (2023). *Theory of mind may have spontaneously emerged in large language models* (arXiv:2302.02083). arXiv. <http://arxiv.org/abs/2302.02083>
- Lange, K., Kühn, S., & Filevich, E. (2015). “Just another tool for online studies” (JATOS): An easy solution for setup and management of web servers supporting online studies. *PLoS One*, 10(6), e0130834. <https://doi.org/10.1371/journal.pone.0130834>
- Lea, R. G., Davis, S. K., Mahoney, B., & Qualter, P. (2019). Does emotional intelligence buffer the effects of acute stress? A systematic review. *Frontiers in Psychology*, 10, 810. <https://doi.org/10.3389/fpsyg.2019.00810>
- Legree, P. J., Psotka, J., Tremble, T., & Bourne, D. R. (2005). Using consensus based measurement to assess emotional intelligence. In *Emotional intelligence: An international handbook* (pp. 155–179). Hogrefe Publishing GmbH.
- Mayer, J. D., Caruso, D. R., & Salovey, P. (2016). The ability model of emotional intelligence: Principles and updates. *Emotion Review*, 8(4), 290–300. <https://doi.org/10.1177/1754073916639667>
- Mayer, J. D., Perkins, D. M., Caruso, D. R., & Salovey, P. (2001). Emotional intelligence and giftedness. *Roeper Review*, 23(3), 131–137. <https://doi.org/10.1080/02783190109554084>
- Mayer, J. D., & Salovey, P. (1995). Emotional intelligence and the construction and regulation of feelings. *Applied and Preventive*

- Psychology*, 4(3), 197–208. [https://doi.org/10.1016/S0962-1849\(05\)80058-7](https://doi.org/10.1016/S0962-1849(05)80058-7)
- Mayer, J. D., Salovey, P., Caruso, D. R., & Sitarenios, G. (2001). Emotional intelligence as a standard intelligence. *Emotion*, 1(3), 232–242. <https://doi.org/10.1037/1528-3542.1.3.232>
- Mayer, J. D., Salovey, P., Caruso, D. R., & Sitarenios, G. (2003). Measuring emotional intelligence with the MSCEIT V2.0. *Emotion*, 3(1), 97–105. <https://doi.org/10.1037/1528-3542.3.1.97>
- McCleskey, J. (2014). Emotional intelligence and leadership: A review of the progress, controversy, and criticism. *International Journal of Organizational Analysis*, 22(1), 76–93. <https://doi.org/10.1108/IJOA-03-2012-0568>
- Moghaddam, S. R., & Honey, C. J. (2023). *Boosting theory-of-mind performance in large language models via prompting*. <https://doi.org/10.48550/ARXIV.2304.11490>
- Nair, V., Schumacher, E., Tso, G., & Kannan, A. (2023). *DERA: Enhancing large language model completions with dialog-enabled resolving agents* (arXiv:2303.17071). arXiv. <http://arxiv.org/abs/2303.17071>
- OpenAI. (2023). *GPT-4 technical report* (arXiv:2303.08774). arXiv. <http://arxiv.org/abs/2303.08774>
- Palmer, B. R., Gignac, G., Manocha, R., & Stough, C. (2005). A psychometric evaluation of the Mayer–Salovey–Caruso emotional intelligence test version 2.0. *Intelligence*, 33(3), 285–305. <https://doi.org/10.1016/j.intell.2004.11.003>
- Peng, B., Alcaide, E., Anthony, Q., Albalak, A., Arcadinho, S., Cao, H., Cheng, X., Chung, M., Grella, M., GV, K. K., He, X., Hou, H., Kazienko, P., Kocon, J., Kong, J., Koptyra, B., Lau, H., Mantri, K. S. I., & Mom, F. (2023). *RWKV: Reinventing RNNs for the Transformer Era*. <https://doi.org/10.48550/ARXIV.2305.13048>
- Salovey, P., & Mayer, J. D. (1990). Emotional intelligence. *Imagination, Cognition and Personality*, 9(3), 185–211. <https://doi.org/10.2190/DUGG-P24E-52WK-6CDG>
- Sap, M., LeBras, R., Fried, D., & Choi, Y. (2023). *Neural theory-of-mind? On the limits of social intelligence in large LMs* (arXiv:2210.13312). arXiv. <http://arxiv.org/abs/2210.13312>
- Shinn, N., Cassano, F., Labash, B., Gopinath, A., Narasimhan, K., & Yao, S. (2023). *Reflexion: Language agents with verbal reinforcement learning*. <https://doi.org/10.48550/ARXIV.2303.11366>
- Sperber, D., & Wilson, D. (2002). Pragmatics, modularity and mind-reading. *Mind & Language*, 17(1-2), 3–23. <https://doi.org/10.1111/1468-0017.00186>
- Taori, R., Gulrajani, I., Zhang, T., Dubois, Y., Li, X., Guestrin, C., Liang, P., & Hashimoto, T. B. (2023b). Stanford Alpaca: An instruction-following LLaMA model. In *GitHub repository*. GitHub. [https://github.com/tatsu-lab/stanford\\_alpaca](https://github.com/tatsu-lab/stanford_alpaca)
- Taori, R., Gulrajani, I., Zhang, T., Dubois, Y., Li, X., Guestrin, C., Liang, P., & Hashimoto, T. B. (2023a). Alpaca: A strong, replicable instruction-following model. *Stanford Center for Research on Foundation Models*, 3(6), 7. <https://crfm.stanford.edu/2023/03/13/alpaca>
- Tian, X., Wang, R., Zhao, Y., Zhen, Z., Song, Y., & Liu, J. (2020). Multi-item discriminability pattern to faces in developmental prosopagnosia reveals distinct mechanisms of face processing. *Cerebral Cortex*, 30(5), 2986–2996. <https://doi.org/10.1093/cercor/bhz289>
- Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Lacroix, T., Rozière, B., Goyal, N., Hambro, E., Azhar, F., Rodriguez, A., Joulin, A., Grave, E., & Lample, G. (2023). *LLaMA: Open and efficient foundation language models* (arXiv:2302.13971). arXiv. <http://arxiv.org/abs/2302.13971>
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is all you need. 31st Conference on Neural Information Processing Systems, Long Beach, CA, USA.
- Wang, X., Zhuang, B., & Wu, Q. (2023). *SwitchGPT: Adapting Large Language Models for Non-Text Outputs*. <https://doi.org/10.48550/ARXIV.2309.07623>
- Warwick, J., & Nettelbeck, T. (2004). Emotional intelligence is ...? *Personality and Individual Differences*, 37(5), 1091–1100. <https://doi.org/10.1016/j.paid.2003.12.003>
- Yang, Z., Li, L., Lin, K., Wang, J., Lin, C.-C., & Liu, Z. (2023). *The Dawn of LMMs: Preliminary Explorations with GPT-4V(ision)*. <https://doi.org/10.48550/ARXIV.2309.17421>
- Zeng, A., Liu, X., Du, Z., Wang, Z., Lai, H., Ding, M., Yang, Z., Xu, Y., Zheng, W., Xia, X., Tam, W. L., Ma, Z., Xue, Y., Zhai, J., Chen, W., Zhang, P., Dong, Y., & Tang, J. (2022). *GLM-130B: An open bilingual pre-trained model* (arXiv:2210.02414). arXiv. <http://arxiv.org/abs/2210.02414>
- Zhao, W. X., Zhou, K., Li, J., Tang, T., Wang, X., Hou, Y., Min, Y., Zhang, B., Zhang, J., Dong, Z., Du, Y., Yang, C., Chen, Y., Chen, Z., Jiang, J., Ren, R., Li, Y., Tang, X., & Liu, Z., ... Wen J.-R. (2023). *A survey of large language models* (arXiv:2303.18223). arXiv. <http://arxiv.org/abs/2303.18223>
- Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E. P., Zhang, H., Gonzalez, J. E., & Stoica, I. (2023). *Judging LLM-as-a-judge with MT-Bench and Chatbot Arena*. <http://arxiv.org/abs/2306.05685>
- Zhu, D., Chen, J., Shen, X., Li, X., & Elhoseiny, M. (2023). *MiniGPT-4: Enhancing Vision-Language Understanding with Advanced Large Language Models*. <https://doi.org/10.48550/ARXIV.2304.10592>